

Feature Maps, Kernels and Regularization

Supervised Learning (COMP0078)

Carlo Ciliberto

AI Centre, Department of Computer Science
University College London

In This Class

1. Beyond Linear Models: Feature Maps
2. Computational Considerations
3. Dual Solution
4. Implicit Feature Maps
5. Kernels
6. Regularization

Part 1: Beyond Linear Models

Recap: Training by fitting some data

The first example of training algorithm:

- Given a dataset of samples, and $(x_i, y_i)_{i=1}^n$
- A loss $\ell(z, y)$ measuring the “error” when predicting z instead of y ,

Goal. Find a function $f : \mathcal{X} \rightarrow \mathcal{Y}$ minimizing


$$\frac{1}{n} \sum_{i=1}^n \ell(f(x), y)$$

Problem. we need¹ to choose a family of candidate models for f !

¹cannot minimize over all f ? **why?**

Recap: Linear Models

Question. What is the simplest form of function we can learn?



For $\mathcal{X} \subset \mathbb{R}^d$ a vector space, we can consider **linear models**

$$f : \mathcal{X} \rightarrow \mathbb{R} \quad \text{such that} \quad f(x) = w^\top x$$

for any $x \in \mathcal{X}$ **input vector** and $w \in \mathbb{R}^d$ the **parameter** vector of f .

Recap: Least Squares Minimization

We then considered $\ell(z, y) = (z - y)^2$ the squared loss and learned

$$\hat{w} = \operatorname{argmin}_{w \in \mathbb{R}^d} \frac{1}{n} \sum_{i=1}^n (w^\top x_i - y_i)^2$$

by:

- Writing the objective in matrix notation $L(w) = \frac{1}{n} \|Xw - y\|^2$
- Setting the derivative $\nabla_w L(w) = 2\frac{1}{n} X^\top Xw - X^\top y = 0$
- solving the corresponding linear system $X^\top X\hat{w} = X^\top y$ yielding

$$\hat{w} = (X^\top X)^{-1} X^\top y$$

What if Linear Models are not Enough?

This is all good and well if x and y are **linearly** related.

But what if they are **not**?

Polynomials. Let us start with the natural non-linear generalization.

Let $\mathcal{X} = \mathbb{R}$ we model $f : \mathcal{X} \rightarrow \mathbb{R}$ as a polynomial of degree $p \geq 0$

$$f(x) = a_0 + a_1x + a_2x^2 + \cdots + a_px^p$$

where $a_0, \dots, a_p \in \mathbb{R}$ are the coefficients.

Polynomial Least Squares

Then, learning a polynomial to fit a training dataset becomes

$$\min_{a_0, \dots, a_p} \frac{1}{n} \sum_{i=1}^n (a_0 + a_1 x_i + a_2 x_i^2 + \dots + a_p x_i^p - y_i)^2$$

(!) **But** the terms x_i^q for $q = 0, \dots, p$ are “just” numbers!

The optimization problem is still **linear** with respect to the parameters a_q

The Polynomial “Feature map”

Define $\phi : \mathcal{X} \rightarrow \mathbb{R}^{p+1}$ such that $\phi(x) = (1, x, x^2, \dots, x^p)^\top$


“Rename” the parameters $w = (w_1, \dots, w_{p+1}) = (a_0, \dots, a_p)$.

Then we can model the polynomial f as


$$f(x) = w^\top \phi(x) \tag{1}$$

Namely a linear model, in the “embedded” space $\phi(\mathcal{X})$


Examples



$$p = 2$$



$$p = 3$$



$$p = 5$$

Polynomial Least Squares (Continued)

Optimization Problem. The training objective becomes

$$\frac{1}{n} \sum_{i=1}^n (w^\top \phi(x_i) - y_i)^2 = \|\Phi w - y\|^2$$

Where $\Phi \in \mathbb{R}^{n \times p}$ is the input matrix, with i -row corresponding to the **embedded** point $\phi(x_i)$ (taking the place of $X \in \mathbb{R}^{n \times d}$).

Solution. Analogously to the linear case, the objective is **minimized** by

$$\hat{w} = (\Phi^\top \Phi)^{-1} \Phi^\top y$$

Same optimization, but we have replaced x with $\phi(x)$!

This construction can be generalized:

- First, to polynomials of degree p over vectors $x \in \mathbb{R}^d$

$$\phi(x) = (x_1^{j_1} x_2^{j_2} \cdots x_d^{j_d})_{j_1, \dots, j_d \in J_p}$$

with indices spanning $J_p = \{(j_1, \dots, j_d) \text{ s.t. } \sum_{i=1}^d j_i = p\}$

- Then, to any $\phi : \mathcal{X} \rightarrow \mathbb{R}^D$ for some “latent” dimension D .

Code-wise

One appealing advantage of this perspective is **modularity**:

- The logic to **learn** the solution \hat{w}

```
import numpy as np

X, y = load_dataset()    # load pairs of input and outputs
Phi = extract_features(X)

hat_w = np.linalg.solve(Phi.T @ Phi, Phi.T @ y)    # "learn"
```

- is separated from that of embedding the input $x \mapsto \phi(x)$

```
# Don't do anything
def extract_linear_features(X):
    return X

# Assume scalar inputs
def extract_quadratic_features(X):
    return np.concatenate((np.ones_like(X), X, np.power(X,2)),
                           axis=1)
```

Part 2: Computations

The Cost of Minimizing the Squared Loss

We have a method for solving the squared loss.

Question. How efficient is it?

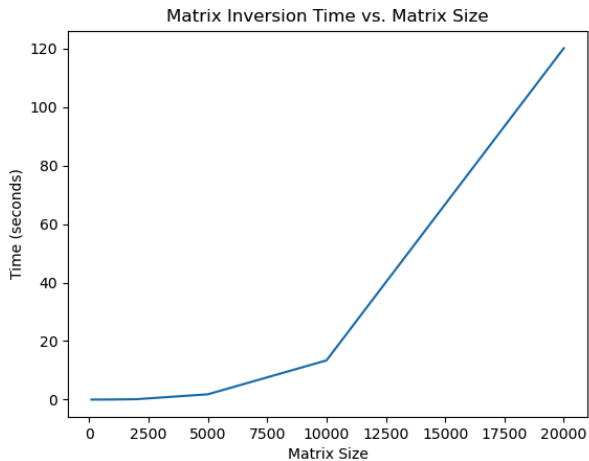
In other words...

how much time (and memory) would it take to obtain \hat{w} ?

Answer. It depends:

- On the machine
- On the training size n
- On the number D of input dimensions.

For example, on my machine...



Costs in Big O notation

Let's try to be more rigorous in asking the question...

Question. What is the cost in **Big O** notation?

Answer:

- $O(D^3 + nD^2)$ in time²,
- $O(D^2 + Dn)$ in memory

²We are referring to standard methods implemented in BLAS/LAPACK

Aside: Big O notation

Informally, $f(x) = O(g(x))$ if it behaves “essentially” like $g(x)$.

Formally³, there exists $C > 0$ and x_0 such that for any $x > x_0$

$$|f(x)| < Cg(x)$$

Note. two functions with “same” big O behavior might differ a lot by:

- The “activating” x_0 (one might behave like $g(x)$ a lot sooner),
- The constant C (one might have a much larger constant).

Hence two methods with same big O, might still be quite different in practice in terms of time and memory.

³Big O notation will be discussed in detail in the TA sessions.
I recommend attending if you are not familiar with it!

When is D too large?

Let assume is fixed, say $n = 1000$. How large can we take D ?

Ideally, as large as possible⁴!

In practice, not too much: cost grows cubically in D .

Examples. Polynomials $f(x)$ of degree p over d variables ($\mathcal{X} = \mathbb{R}^d$):

- $d = 1, p = 3 \implies D = 4$
- $d = 3, p = 5 \implies D = 56$
- $d = 5, p = 10 \implies D = 3003$
- $d = 10, p = 15 \implies D = 184756$

⁴Well... not really because of overfitting (we will get back to this point), but we would not like computations to be our bottleneck!

A Closer Look

Where is the cost $O(D^3 + nD^2)$ coming from?

$$\hat{w} = (\underbrace{\Phi^T \Phi}_{\text{Multiplication}}) \underbrace{-1}_{\text{Inversion}} \Phi^T y$$

$(D \times n) \cdot (n \times D) \rightarrow O(nD^2)$

$D \times D \rightarrow O(D^3)$

Is there any way we could estimate \hat{w} faster?

Part 3: Dual Solution

Alternative Formulation for Squared Loss Minimization

We will manipulate the solution

$$\hat{w} = (\Phi^\top \Phi)^{-1} \Phi^\top y$$

to derive:

- an **equivalent** characterization
- that can be **faster** in some regimes.

Our Strategy

There are many ways to approach this problem.

We will use a strategy that makes significant use of **linear algebra**.

This will familiarize us with tools that we will use also in the future.

Singular Value Decomposition

We will start by considering the **Singular Value Decomposition (SVD)**

$$\Phi = U\Sigma V^\top$$

Where:

- $r = \min(n, D)$
- $U \in \mathbb{R}^{n \times r}$ and $V \in \mathbb{R}^{D \times r}$ have **orthonormal** columns

$$U^\top U = I \quad \text{and} \quad V^\top V = I$$

- $\Sigma \in \mathbb{R}^{r \times r}$ is a **diagonal** matrix with only non-negative entries.
- The diagonal vector $\sigma = \text{diag}(\Sigma) \in \mathbb{R}^r$ is the **spectrum** of Σ
- The number of non-zero entries in σ is the rank of Φ .

The SVD and other useful prerequisites will be discussed also in the TA sessions. Attendance is recommended.

SVD and Squared Loss Minimizer

(!) For Simplicity. Assume Φ “full rank”: $\text{rank}(\Phi) = r$.

Let us plug in the SVD of $\Phi = U\Sigma V^\top$ into \hat{w}

$$\hat{w} = (\Phi^\top \Phi)^{-1} \Phi^\top y = (V\Sigma U^\top U\Sigma V^\top)^{-1} V\Sigma U^\top y$$

- Since $U^\top U = I$, we have $V\Sigma U^\top U\Sigma V^\top = V\Sigma^2 V^\top$
- Since $V^\top V = I$, we have $(V\Sigma^2 V^\top)(V\Sigma^{-2} V^\top) = I$. Hence

$$(V\Sigma^2 V^\top)^{-1} = V\Sigma^{-2} V^\top$$

- Combining the two

$$\hat{w} = (V\Sigma^{-2} V^\top) V\Sigma U^\top y = V\Sigma^{-1} U^\top y$$

SVD and Squared Loss Minimizer (Continued)

We can now “go back” by making the pairs $U^\top U$ and $V^\top V$ “appear”

$$\begin{aligned}\hat{w} &= V\Sigma^{-1}U^\top y \\ &= V\Sigma\Sigma^{-2}U^\top y \\ &= V\Sigma U^\top U\Sigma^{-2}U^\top y \\ &= \underbrace{(V\Sigma U^\top)}_{\Phi^\top} (U\Sigma^{-2}U^\top) y \\ &= \Phi^\top \underbrace{(U\Sigma^2 U^\top)^{-1}}_{\Phi\Phi^\top} y \\ &= \Phi^\top (\Phi\Phi^\top)^{-1} y\end{aligned}$$

SVD and Squared Loss Minimizer (Continued)

We can now “go back” by making the pairs $U^\top U$ and $V^\top V$ “appear”

$$\begin{aligned}\hat{w} &= V\Sigma^{-1}U^\top y \\ &= V\Sigma\Sigma^{-2}U^\top y \\ &= V\Sigma U^\top U\Sigma^{-2}U^\top y \\ &= \underbrace{(V\Sigma U^\top)}_{\Phi^\top} (U\Sigma^{-2}U^\top) y \\ &= \Phi^\top \underbrace{(U\Sigma^2 U^\top)^{-1}}_{\Phi\Phi^\top} y \\ &= \Phi^\top (\Phi\Phi^\top)^{-1} y\end{aligned}$$

(!) **Note** in both this and the previous slide we have been a bit sloppy about matrix inversion. We will get back to this soon, to make our derivation more rigorous.

Dual vs “Primal” formulation

After a bit of work we found out that \hat{w} can be written as

$$\hat{w} = \Phi(\Phi^{-1}\Phi)^{-1}y$$

...and... so what?

Dual vs “Primal” formulation

After a bit of work we found out that \hat{w} can be written as

$$\hat{w} = \Phi(\Phi^{-1}\Phi)^{-1}y$$

...and... so what?

At a superficial look, this formula is not different from the “primal”

$$\hat{w} = (\Phi^{\top}\Phi)^{-1}\Phi^{\top}y$$

We just “swapped” Φ^{\top} and $(\Phi^{\top}\Phi)^{-1}$...

...or did we?

The Benefits of the Dual Formulation

$$(\Phi^\top \Phi)^{-1} \Phi^\top y \quad \text{Vs} \quad \Phi (\Phi^{-1} \Phi)^{-1} y$$

The matrices inverted have different dimensions!

- $\Phi^\top \Phi$ is $D \times D$
- $\Phi \Phi^\top$ is $n \times n$.

Therefore the computational costs for the two formulations are:

- **Primal:** $O(D^3 + nD^2)$
- **Dual:** $O(n^3 + Dn^2)$

**Therefore in the Dual formulation the dominating factor is $n!$
(wrt which cost grows cubically)**

The Benefits of the Dual Formulation (Continued)

Of course, this does not mean that the dual formulation is “better”

It simply shows that **when** $N < D$ the dual formulation is faster.

This is appealing, however, since we have seen (e.g. the polynomials) that the feature space can **grow** pretty quickly wrt the problem parameters!

In particular, it will be critical when we will consider feature space that are **infinite dimensional!**

Dotting our “i”s and crossing our “t”s: (Pseudo)inversion

We have been sloppy in our use of the notation $(\Phi^\top \Phi)^{-1}$ and $(\Phi \Phi^\top)^{-1}$.

If $n \neq D$ however, **at least one** of the two matrices will not be invertible.

More more generally, if $\text{rank}(\Phi) < \min(n, D)$, **neither** will.

The correct quantity that we should have used is the **Pseudoinverse**.

Def. The pseudoinverse⁵ of $\Phi \in \mathbb{R}^{n \times D}$ is a matrix $\Phi^\dagger \in \mathbb{R}^{D \times n}$ such that

- Behaves like a “weak inverse”:

$$\Phi \Phi^\dagger \Phi = \Phi \quad \text{and} \quad \Phi^\dagger \Phi \Phi^\dagger = \Phi^\dagger,$$

- Both $\Phi \Phi^\dagger$ and $\Phi^\dagger \Phi$ are Symmetric.

⁵can be generalized to Complex matrices.

The SVD of the pseudoinverse of $\Phi = U\Sigma V^\top$ is

$$\Phi^\dagger = V\Sigma^\dagger U^\top$$

Where Σ^\dagger is the diagonal matrix with entries

$$(\Sigma^\dagger)_{ij} = \begin{cases} \frac{1}{\Sigma_{ij}} & \text{if } \Sigma_{ij} \neq 0 \\ 0 & \text{otherwise} \end{cases} \quad (2)$$

Dotting our “i”s and crossing our “t”s: (Pseudo)inversion

We can therefore retrace our derivation of the dual formulation⁶ to obtain

$$\hat{w} = \underbrace{(\Phi^\top \Phi)^\dagger \Phi^\top}_{\textit{Primal}} y = \Phi^\top \underbrace{(\Phi \Phi^\top)^\dagger}_{\textit{Dual}} y$$

which is now well-defined.

⁶Reccomendation: do that as an exercise!

In practice... Pseudoinverse calculation is often numerically unstable⁷ (we will get back to this later).

It is recommended to use other methods. For example **linear solvers**⁸

Primal

```
hat_w = np.linalg.solve(Phi.T @ Phi, Phi.T @ y)
```

Dual

```
alpha = np.linalg.solve(Phi @ Phi.T, y)      # dual coefficients  
w_hat = Phi.T @ alpha
```

⁷meaning you might get the wrong \hat{w} out!

⁸Namely routine for solving problem of the form $Ax = b$, which is indeed $x = A^\dagger b$

Part 4: Implicit Feature Maps

A Deeper Look into the Dual Formulation

The code snippet for the **dual** formulation

```
alpha = np.linalg.solve(Phi @ Phi.T, y)    # dual coefficients  
w_hat = Phi.T @ alpha
```

highlights an interesting structure.

Denote $K = \Phi\Phi^\top$, then the solution \hat{w} is

$$\hat{w} = \Phi^\top \alpha \quad \text{where} \quad \alpha = K^{-1}y$$

In some sense, we have **decoupled**:

- The “training” (the computation intensive inversion of α).
- The evaluation $\hat{f}(x) = \phi(x)^\top \hat{w} = \phi(x)\Phi^\top \alpha$.

(?) Why is this interesting?

Inner Products

We have two interesting observations:

1. $K_{ij} = (\Phi\Phi^\top)_{ij} = \phi(x_i)^\top \phi(x_j),$
2. Evaluating the learned function \hat{f} in a new point x yields

$$\hat{f}(x) = \phi(x)^\top \hat{w} = \phi(x) \Phi^\top \alpha = v(x)^\top \alpha$$

where we have defined the **evaluation vector** $v(x) \in \mathbb{R}^n$

$$v(x) = (\phi(x)^\top \phi(x_1), \dots, \phi(x)^\top \phi(x_n))^\top$$

Inner Products

We have two interesting observations:

1. $K_{ij} = (\Phi\Phi^\top)_{ij} = \phi(x_i)^\top \phi(x_j)$,
2. Evaluating the learned function \hat{f} in a new point x yields

$$\hat{f}(x) = \phi(x)^\top \hat{w} = \phi(x) \Phi^\top \alpha = v(x)^\top \alpha$$

where we have defined the **evaluation vector** $v(x) \in \mathbb{R}^n$

$$v(x) = (\phi(x)^\top \phi(x_1), \dots, \phi(x)^\top \phi(x_n))^\top$$

All the operations required to train and evaluate \hat{w} depend only on the inner products $\phi(x)^\top \phi(x')$!

Are Inner Products Better

Question. Are there feature maps ϕ for which we can compute $\phi(x)^\top \phi(x')$ “directly” from x and x' ...

...without explicitly needing the embedded vectors?

Are Inner Products Better

Question. Are there feature maps ϕ for which we can compute $\phi(x)^\top \phi(x')$ “directly” from x and x' ...

...without explicitly needing the embedded vectors?

Yes!

Example. Consider $x \in \mathcal{X} \subset \mathbb{R}^2$ and $\phi : \mathbb{R}^2 \rightarrow \mathbb{R}^3$, hence $D = 3$

$$\phi(x) = (x_1^2, \sqrt{2}x_1x_2, x_2^2)^\top$$

In principle $\phi(x)^\top \phi(x')$ requires the sum of $D = 3$ products. However...

$$\begin{aligned}\phi(x)^\top \phi(x') &= \underbrace{x_1^2 x_1'^2}_{a^2} + 2 \underbrace{x_1 x_1'}_a \underbrace{x_2 x_2'}_b + \underbrace{x_2^2 x_2'^2}_{b^2} \\ &= (a + b)^2\end{aligned}$$

...one sum and one product are **sufficient!**

Spaces of Polynomials and Inner Product

The example generalizes to polynomials of degree p in dimension d :

Proposition. For any $p, d \in \mathbb{N}$ there exist coefficients $C_{j_1, \dots, j_d} \in \mathbb{R}$ and a feature map $\phi : \mathbb{R}^d \rightarrow \mathbb{R}^D$

$$\phi(x) = \left(C_{j_1, \dots, j_d} x_1^{j_1} \cdots x_d^{j_d} \right)_{\sum_i j_i \leq p}$$

such that

$$\phi(x)^\top \phi(x') = (x^\top x' + 1)^p$$

Space of Polynomials and Inner Product

A few observations:

- The C_{j_1, \dots, j_d} can be obtained by expanding the product $(x^\top x' + 1)^p$
- $D = O(p^d)$ (D = the number of addends in the expansion above).
- Evaluating $(x^\top x' + 1)^p$ requires “only” $O(d)$ operations.

⁹Another, non-trivial advantage: we **don't need** to calculate the C_{j_1, \dots, j_d} explicitly!

Space of Polynomials and Inner Product

A few observations:

- The C_{j_1, \dots, j_d} can be obtained by expanding the product $(x^\top x' + 1)^p$
- $D = O(p^d)$ (D = the number of addends in the expansion above).
- Evaluating $(x^\top x' + 1)^p$ requires “only” $O(d)$ operations.

Therefore, the computational costs are:

- **Primal.** $O(p^{3d} + np^{2d})$
- **Dual.** $O(n^3 + p^d n^2)$
- **Inner Product + Dual.** $O(n^3 + dn^2)$

Therefore Dual formulation + Using only inner product provides a huge computational⁹ advantage!

⁹Another, non-trivial advantage: we **don't need** to calculate the C_{j_1, \dots, j_d} explicitly!

Wishlist. More generally, we would like to find feature maps ϕ for which there exists a function $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ such that

$$k(x, x') = \phi(x)^\top \phi(x')$$

and the calculation **does not** require explicit knowledge of $\phi(x)$

Then. If evaluating $k(x, x')$ can be done in $O(\kappa)$ operations...
...it follows that learning \hat{w} requires

$$O(n^3 + \kappa n^2)$$

with the **Inner Product + Dual** formulations¹⁰

¹⁰Recall that you will still need $O(n)$ operations to evaluate $f(x)$!

Inverting the Question

Now, consider our current question...

Can we find a feature map ϕ whose associated inner product k that can be evaluated without knowledge of $\phi(x)$?

...and let us turn it on its head by asking ourselves:

When does a function $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ correspond to the inner product of some feature map ϕ ?

Why might this be interesting?

Searching for an Inner Product

Intuition. Inner products are a measure of similarity (actually correlation) between two objects. If $\phi(x)^\top \phi(x')$ is:

- **Positive and large.** Then x and x' are “similar” according to ϕ .
- **Close to zero.** Then x and x' are uncorrelated (not very similar)
- **Negative and large.** Then x and x' are oppositely correlated.

Sometimes it might be easier¹¹ to describe a desirable notion of similarity/correlation as a function $k(x, x')$ explicitly, rather than implicitly by designing a feature map ϕ .

BUT then, the question is whether such function is an inner product!

¹¹another advantage: by designing k , we could choose one that is fast to compute!

So our next question will be

Can we find a criterion to determine when a function $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ corresponds to the inner product of some feature map ϕ ?

Better yet, we look for some constructive tools

Can we find some rules that allow us to build functions $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ that guarantee they correspond to the inner product of some feature map ϕ ?

Part 5: Kernels

(Implicitly) Recovering Inner Products

Goal. Find an **alternative** characterization for an inner product function $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ that **does not require** explicit knowledge of ϕ .

How? Let's look at some properties of inner product functions.

Observation. Let $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$. For any set of points $(x_i)_{i=1}^n$, let $K \in \mathbb{R}^{n \times n}$ be the corresponding “Gram” matrix with entries

$$K_{ij} = k(x_i, x_j).$$

Then.¹² k is an inner product $\implies K$ is **Positive Semidefinite**¹³ ($K \succeq 0$).

¹²Exercise! (Actually... we already proved some slides ago)

¹³This holds for any Gram matrix K (namely any $n \in \mathbb{N}$ and any set $(x_i)_{i=1}^n$)

Positive Semidefinite Gram Matrices

Recall. A matrix $K \in \mathbb{R}^{n \times n}$ is **Positive Semidefinite (PSD)** if it is symmetric and $v^\top K v \geq 0$ for any $v \in \mathbb{R}^n$

Equivalently K is PSD \iff it's SVD is of the form¹⁴ $K = U \Sigma U^\top$.

Is it important for K to be PSD?

¹⁴Note: here $U = V$ and recall that Σ has only non-negative entries

(Implicitly) Recovering Inner Products

Question. Is the converse “ k is an inner product $\iff K \succeq 0$ ” true?

Maybe. But is it important?

Strictly speaking, **it is not:** we can still define the estimator

$$\hat{f}(x) = v(x)^\top \alpha \quad \text{with} \quad \alpha = K^\dagger y$$

and $v(x) = (k(x, x_1), \dots, k(x, x_n))^\top$ the “evaluation” vector.

¹⁵But if all this has not scared you, have a look at *reproducing kernel Krein spaces*

(Implicitly) Recovering Inner Products

Question. Is the converse “ k is an inner product $\iff K \succeq 0$ ” true?

Maybe. But is it important?

Strictly speaking, **it is not:** we can still define the estimator

$$\hat{f}(x) = v(x)^\top \alpha \quad \text{with} \quad \alpha = K^\dagger y$$

and $v(x) = (k(x, x_1), \dots, k(x, x_n))^\top$ the “evaluation” vector.

However, this approach introduces a series of:

- **Numerical** issues in practice,
- **Methodological** issues algorithmically,
- **Statistical** issue theoretically.

In other words, it's a mess¹⁵!

¹⁵But if all this has not scared you, have a look at *reproducing kernel Krein spaces*

Positive Definite Kernels

So, let's assume that we **want** k to be PSD.

Def. We say that $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ is a *positive definite kernel* if for any dataset $(x_i)_{i=1}^n$ the corresponding Gram matrix is PSD.

So... is this enough for k to be an inner product?

If it Quacks Like a Duck...

Theorem. k is an inner product $\iff k$ is a positive definite kernel.

Yay!

In Depth. What can be shown is that there exist

- A **Hilbert** space \mathcal{H} with **inner product** $\langle \cdot, \cdot \rangle_{\mathcal{H}}$
- A **feature map** $\phi : \mathcal{X} \rightarrow \mathcal{H}$

Such that

$$k(x, x') = \langle \phi(x), \phi(x') \rangle_{\mathcal{H}} \quad \forall x, x' \in \mathcal{X}$$

Which is precisely what we wanted!

Aside: Hilbert Spaces

Recall. A **Hilbert** space is:

- A vector space (in our case we assume it to be real)
- Equipped with an **inner product** $\langle \cdot, \cdot \rangle_{\mathcal{H}}$
- ...and the corresponding **norm** $\| \cdot \|_{\mathcal{H}}$ such that $\|x\|_{\mathcal{H}}^2 = \langle x, x \rangle_{\mathcal{H}}$
- That is **complete**¹⁶ with respect to $\| \cdot \|_{\mathcal{H}}$.

Recall. An **inner product** $\langle \cdot, \cdot \rangle$ over a (real) vector space \mathcal{V} is a function $\langle \cdot, \cdot \rangle : \mathcal{V} \times \mathcal{V} \rightarrow \mathbb{R}$ such that $\forall v, v', v'' \in \mathcal{V}$ and $\alpha, \beta \in \mathbb{R}$:

- **Symmetric** $\langle v, v' \rangle = \langle v', v \rangle$
- **(Bi)linear** $\langle \alpha v + \beta v', v'' \rangle = \alpha \langle v, v'' \rangle + \beta \langle v', v'' \rangle$
- **Positive definite** $\langle v, v \rangle > 0$

¹⁶that is, every Cauchy sequence admits a limit in \mathcal{H} .

Infinite Dimensional Hilbert Spaces

Note. when \mathcal{H} is finite of dimension D , then $\mathcal{H} \cong \mathbb{R}^D$.

So far, we have only considered spaces of this form, but nothing prevents us to consider **infinite dimensional spaces**.

Example (The Gaussian Kernel) Let $\sigma > 0$. Then the function

$$k(x, z) = e^{-\|x-z\|^2/2\sigma}$$

is a kernel associated to an infinite Hilbert space (we will show this).

- Using **primal** or (naive) **dual** would be impossible ($D = +\infty$)!
- With **Kernel + Dual** it takes $O(n^3 + dn^2)$ ($k(x, z)$ requires $O(d)$).

The Gaussian Kernel Feature Map

Let's find a feature map for the Gaussian kernel.

For simplicity let $\sigma = 1$. Also, we will show¹⁷ this for $\mathcal{X} = \mathbb{R}$.

- First, observe that by expanding the square

$$k(x, z) = e^{-(x-z)^2/2} = e^{-(x^2/2 - xz + z^2/2)} = e^{-x^2/2} e^{xz} e^{-z^2/2}$$

- Then, consider the Taylor series of the exponential for the mixed term

$$e^{xz} = \sum_{m=1}^{+\infty} \frac{(xz)^m}{m!} = \sum_{m=1}^{+\infty} \frac{x^m}{\sqrt{m!}} \frac{z^m}{\sqrt{m!}}$$

¹⁷The derivation becomes more involved for \mathbb{R}^d with $d > 1$

The Gaussian Kernel Feature Map (Continued)

Let ℓ^2 denote the space of all (real) square-summable sequences¹⁸

Recall. ℓ^2 is a Hilbert space with $\langle (a_m)_{m=1}^{+\infty}, (b_m)_{m=1}^{+\infty} \rangle = \sum_{m=1}^{+\infty} a_m b_m$.

Define¹⁹ $\phi : \mathbb{R} \rightarrow \ell^2$ the map such that $\forall x \in \mathbb{R}$

$$\phi(x) = \left(e^{-x^2/2} \frac{x^m}{\sqrt{m}} \right)_{m=1}^{+\infty}$$

Then,

$$\langle \phi(x), \phi(z) \rangle_{\ell^2} = \sum_{m=1}^{+\infty} e^{-x^2/2} \frac{x^m}{\sqrt{m}} \frac{z^m}{\sqrt{m}} e^{-z^2/2} = e^{-(x-z)^2/2}$$

¹⁸namely sequences $(a_m)_{m=1}^{+\infty}$ such that $\sum_{m=1}^{+\infty} a_m^2 < +\infty$.

¹⁹Is ϕ well defined (i.e. $\phi(x) \in \ell^2$?). **Exercise!**

The Gaussian Kernel Feature Map (Continued)

We proved that the Gaussian kernel is indeed a kernel by **explicitly** finding \mathcal{H} and a feature map ϕ .

This was fun, but doing it **for every kernel** gets old pretty quickly.

As already stressed many times, we **don't need** to know the feature map associated to a kernel. So...

Is there a way to show that a function is a kernel without going the “feature map way”?

Theorem. Let $k_1, k_2 : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ be two positive definite kernels, then

- αk_1 is a kernel for any $\alpha \geq 0$,
- **Sum.** $k_1 + k_2$ is a kernel,
- **Point-wise product** $k(x, x') = k_1(x, x')k_2(x, x')$ is a kernel,
- **Composition** $k_1(\varphi(z), \varphi(z'))$ is a kernel for any $\varphi : \mathcal{Z} \rightarrow \mathcal{X}$

Let us look at a few examples...

Polynomial Kernel. The function

$$k(x, x') = (x^\top x' + \alpha)^p$$

with $\alpha \geq 0$ is a kernel because:

- $k_1(x, x') = x^\top x'$ is a kernel,
- a constant function $k_2(x, x') = \alpha$ is a kernel (exercise?),
- $k_3 = k_1 + k_2$ is a kernel,
- $k = k_3^p$ is a power of a kernel (k_3 times itself p times).

Examples of Kernels: Polynomial of Kernels

Polynomials OF Kernels The function

$$k(x, x') = \sum_{m=1}^p \alpha_m \tilde{k}(x, x')^m$$

Is a kernel (sum of powers of kernels).

Series of Kernels. We could send $m \rightarrow +\infty$ in the above construction.
If the series converges **then it is a kernel!**

For example, let $\alpha_m = 1/m!$ and $\tilde{k}(x, x') = (x^\top x')^m$ the polynomial kernel, we obtain

$$k(x, x') = \sum_{m=1}^{+\infty} \frac{(x^\top x')^m}{m!} = e^{xx'}$$

the exponential kernel.

The Gaussian Kernel (Again)

We can use the rules above to show that $k(x, x') = e^{-\|x-x'\|^2}$ is a kernel.

Expanding (again) the square we have

$$k(x, x') = e^{-(\|x\|^2/2 - xx' + \|x'\|^2/2)}$$

We can write the above equation as the point-wise product

$$k(x, x') = \underbrace{e^{-\|x\|^2} e^{-\|x'\|^2}}_{k_1(x, x')} \underbrace{e^{xx'}}_{k_2(x, x')}$$

Where,

- $k_1(x, x') = \phi_1(x)\phi(x')$ is a kernel with **feature map** $\phi(x) = e^{-\|x\|^2}$
- k_2 is the **exponential kernel**.

Hence k is a kernel.

Going Deeper with Kernels

We barely scratched the **surface** of what can be studied about kernels.

Given the scope of this module, we will limit ourselves to the **algorithmic perspective** on kernels and corresponding Hilbert spaces that we introduced in this class.

However, if you are interested in this area, I would strongly recommend **COMP0083 Advanced Topics in Machine Learning (ATML)**, which offers a nice complement to this module with topics on:

- Reproducing Kernel Hilbert Spaces
- Convex optimization

Part 6: Regularization

The Problem with Large Feature Spaces

Kernels allow to **efficiently** work with **large (or infinite!)** feature spaces!

However, there is a potential problem²⁰...

The number n of training points is:

- **Finite** (can't really have an infinite dataset, can we?)
- Potentially **Smaller** than the feature space dimension D
(Think about polynomial or Gaussian kernels).


²⁰isn't there always one?

Fitting Linear Models

A linear $f : \mathbb{R}^D \rightarrow \mathbb{R}$ is **uniquely** characterized by $n = D$ points.

(for f affine you need $n = D + 1$)

If $n < D$ you will have **more than one** possible functions.

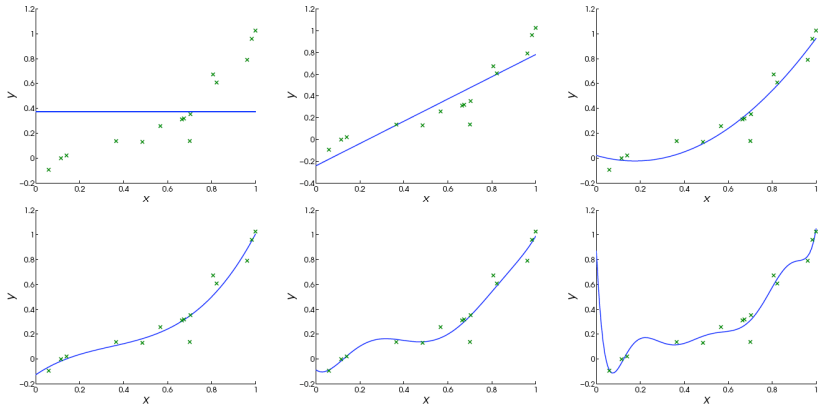


This is because $\hat{w} = (\Phi^\top \Phi)^{-1} \Phi^\top y$ corresponds to **solving** $\Phi \hat{w} = y$

Namely linear system of:

- n **equations** (one per point),
- D **variables**

Over Fitting your Points



Example: same number of points, polynomials of degree $p = 0$ to 5.

How Does Overfitting Arise?

Mechanically, we can see this issue arising in $\hat{w} = (\Phi^\top \Phi)^{-1} \Phi^\top y$ since

$$n < D \implies \Phi^\top \Phi \text{ not invertible}$$

However, we saw that the lack of invertibility it is not a problem since:

- We can use the **dual** $\hat{w} = \Phi^\top (\Phi \Phi^\top)^{-1} y$, or better
- Use the **pseudoinverse**, rather than the inverse, of $\Phi \Phi^\top$.

However, the Pseudoinverse is not a panacea.

Instability of the Pseudoinverse on Small Principal Directions

Let $\Phi = U\Sigma V^\top$, with rank $r = n$ and smallest singular value $\sigma_n = 10^{-p}$.

We have shown that the squared loss minimizer is

$$\hat{w} = (\Phi^\top \Phi)^\dagger \Phi^\top y = V\Sigma^\dagger U^\top y$$

Let $x = v_n$ be the n -th column of V . Let $x' = (1 + \epsilon)x$ and $\epsilon \in \mathbb{R}$. Then

$$|\hat{f}(x) - \hat{f}(x')| = |\epsilon| |v_n^\top V\Sigma^\dagger U^\top y| = \frac{|\epsilon|}{\sigma_n} |u_n^\top y|$$

Since $\sigma_n = 10^{-p}$, this implies that two points x and x' that are close ϵ to each other, will be predicted to have outputs that are

$$10^p \epsilon |u_n^\top y|$$

Namely, small “input” perturbations along v_n will be inflated by $\times 10^p$

Monkey-Patching the Pseudoinverse?

If the problem are small singular values, can we just get rid of them?

Idea. Choose a threshold $\lambda > 0$ and set to zero any $\Sigma_{ii} \leq \lambda$

Then, we would have the solution

$$\hat{w}_\lambda = VP_\lambda(\Sigma)^\dagger U^\top y$$

Where²¹ $P_\lambda : \mathbb{R}^{n \times n} \rightarrow \mathbb{R}^{n \times n}$

$$\left(P_\lambda(\Sigma)\right)_{ii} = \begin{cases} \frac{1}{\Sigma_{ii}} & \text{if } \Sigma_{ii} > \lambda \\ 0 & \text{otherwise} \end{cases}$$

²¹More rigorously, P_λ sends diagonal matrices into diagonal matrices.

Regularization(s)

Yes we can! $\hat{w}_\lambda = VP_\lambda(\Sigma)^\dagger U^\top y$ is one example of **regularization**.

We have seen another example in the previous lecture:

$$\hat{w}_\lambda = (\Phi^\top \Phi + \lambda I)^{-1} \Phi^\top y$$

that has also an interpretation as the solution of

$$\hat{w}_\lambda = \underset{w \in \mathbb{R}^D}{\operatorname{argmin}} \frac{1}{n} \sum_{i=1}^n (w^\top \phi(x_i) - y_i)^2 + \lambda \|w\|^2$$

Also this solution “dampens” small singular values since

$$\hat{w}_\lambda = V(\Sigma^2 + \lambda I)^{-1} \Sigma U^\top y$$

Thus, the smallest singular value is always larger than $1/\lambda$.

Questions...

By modifying \hat{w} into \hat{w}_λ we are not fitting anymore our data...

$$\frac{1}{n} \sum_{i=1}^n (\hat{w}^\top \phi(x_i) - y_i)^2 \leq \frac{1}{n} \sum_{i=1}^n (\hat{w}_\lambda^\top \phi(x_i) - y_i)^2$$

Is this a problem? Maybe not.

Recall. we actually care about the expected risk

$$\mathcal{E}(\hat{w}_\lambda) = \mathbb{E}_{(x,y)}[(\hat{w}_\lambda^\top \phi(x) - y)^2]$$

Therefore (in principle at least) it might even happen that

$$\mathcal{E}(\hat{w}_\lambda) \leq \mathcal{E}(\hat{w})$$

Bias and Variance

Regularization has some **pros** and **cons**. **Recall:** let S a training set

$$\mathbb{E}_S[\mathcal{E}(\hat{w}_\lambda) - \mathcal{E}(f_*)] = \textit{Bias}(\hat{w}_\lambda, f_*) + \textit{Variance}(\hat{w}_\lambda)$$

- **Bias.** “Average” distance from the ground truth

$$\textit{Bias}(\hat{w}_\lambda, f_*) = \mathbb{E}_x[(f_*(x) - \hat{w}_\lambda^\top \phi(x))^2]$$

Hence, if $f_*(x) = w^\top \phi(x)$, the larger λ the **further away** our predictions will be... **BAD!**

- **Variance.** Variance of models' output wrt training set S

$$\textit{Variance}(\hat{w}_\lambda) = \mathbb{E}_{x,S}[|(\hat{w}_\lambda - \mathbb{E}_s[\hat{w}_\lambda])^\top \phi(x)|^2]$$

The larger λ the smaller the variance (**why?**)... **GOOD!**

Model Selection

Ok, so in theory we might have an advantage...

...but in practice what should we do?

Model Selection. Choose the best λ according \hat{w}_λ 's performance:

- **Hold-out.** Split training set $S = S_{tr} \cup S_{val}$
 - Choose a few candidate $\lambda_1, \dots, \lambda_m$
 - Train \hat{w}_λ on S_{tr} for each λ_j
 - Evaluate each \hat{w}_λ on S_{va}
 - Choose \hat{w}_{λ_*} the best performing $\lambda_* = \lambda_{j_*}$
- **K-fold Cross-validation.** Consider K train-val splits of S .
 - Perform hold-out for each of the splits.
 - Return λ_* with the best w_{λ_*} on average.
- ...

Wrapping up

In this class:

- We have introduced feature maps to go beyond linear models,
- Leverage the Dual formulation when $n < D$
- Observed how inner products might be performed “implicitly”,
- Introduced rules to build kernels,
- Discusses how to regularize learning problems. =

Next class:

we will focus on classification problems and consider a method tailored for these setting: Support Vector Machines.